

## Importing Modules

```
In [1]: import sys
sys.executable
Out[1]: 'C:\Users\hardi\anaconda3\python.exe'

In [2]: import librosa
import librosa.display

In [3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from IPython.display import Audio
import warnings
warnings.filterwarnings('ignore')

In [4]: import os
path = ''
labels = []
dataset_location = 'C:\Users\hardi\Downloads\archive\TESS Toronto emotional speech set data'
for dirname, _, filenames in os.walk(dataset_location):
    for filename in filenames:
        file_path = os.path.join(dirname, filename)
        label = filename.split('.')[1]
        paths.append(file_path)
        labels.append(label.lower())
print('Dataset is loaded')
Dataset is loaded

In [5]: paths[5]

Out[5]: ['C:\Users\hardi\Downloads\archive\TESS Toronto emotional speech set data\DAF_angry\DAF_angry.wav',
'C:\Users\hardi\Downloads\archive\TESS Toronto emotional speech set data\DAF_angry\DAF_angry.wav',
'C:\Users\hardi\Downloads\archive\TESS Toronto emotional speech set data\DAF_angry\DAF_base_angry.wav',
'C:\Users\hardi\Downloads\archive\TESS Toronto emotional speech set data\DAF_angry\DAF_base_angry.wav',
'C:\Users\hardi\Downloads\archive\TESS Toronto emotional speech set data\DAF_angry\DAF_base_angry.wav']

In [6]: labels[5]

Out[6]: ['angry', 'angry', 'angry', 'angry', 'angry']

In [7]: df = pd.DataFrame()
df['speech'] = paths
df['label'] = labels
df.head()
```

```
Out[7]:
```

	speech	label
0	C:\Users\hardi\Downloads\archive\TESS Toronto ...	angry
1	C:\Users\hardi\Downloads\archive\TESS Toronto ...	angry
2	C:\Users\hardi\Downloads\archive\TESS Toronto ...	angry
3	C:\Users\hardi\Downloads\archive\TESS Toronto ...	angry
4	C:\Users\hardi\Downloads\archive\TESS Toronto ...	angry

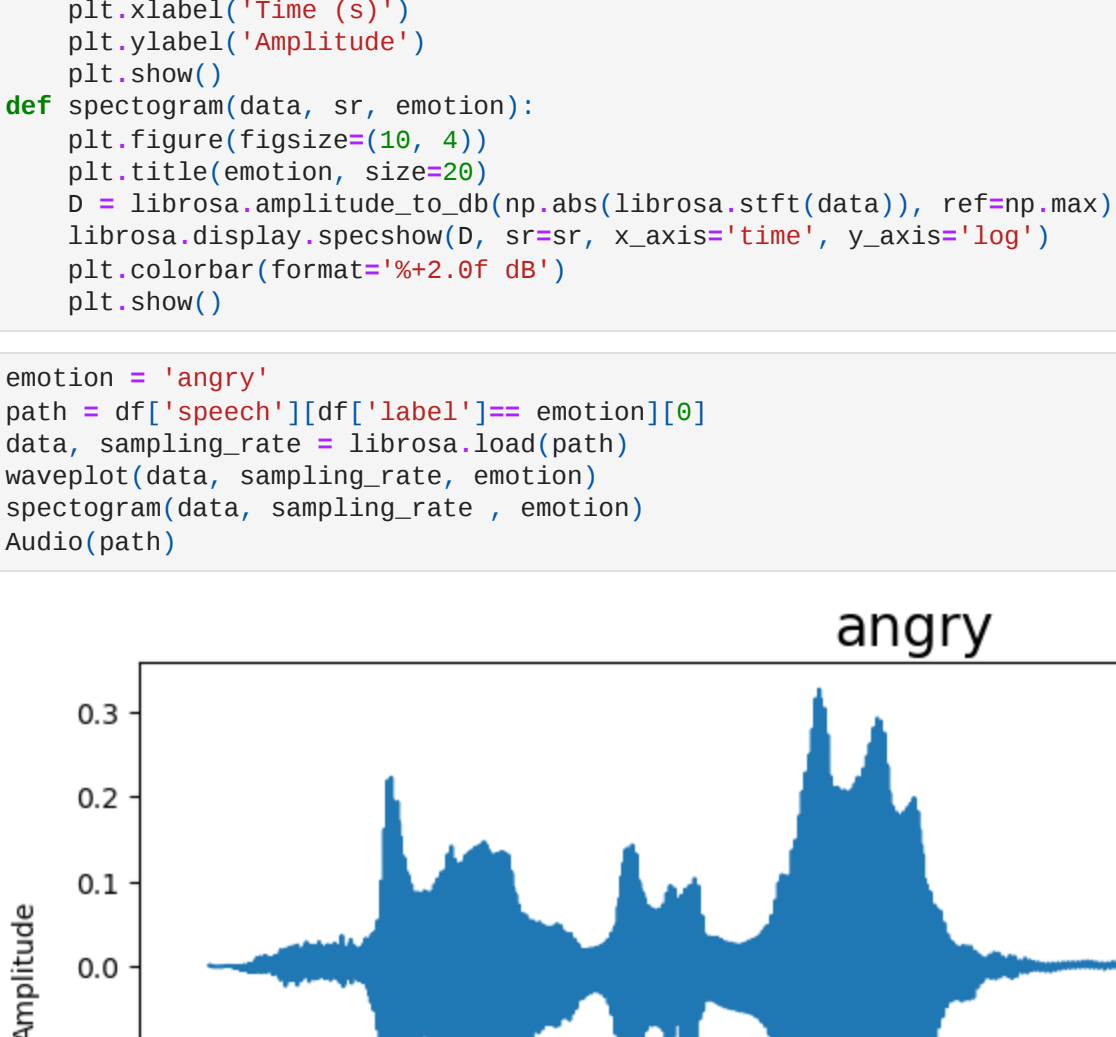
```
In [8]: df['label'].value_counts()
Out[8]:
```

label	count
angry	800
disgust	800
fear	800
happy	800
neutral	800
ps	800
sad	800

Name: label, dtype: int64

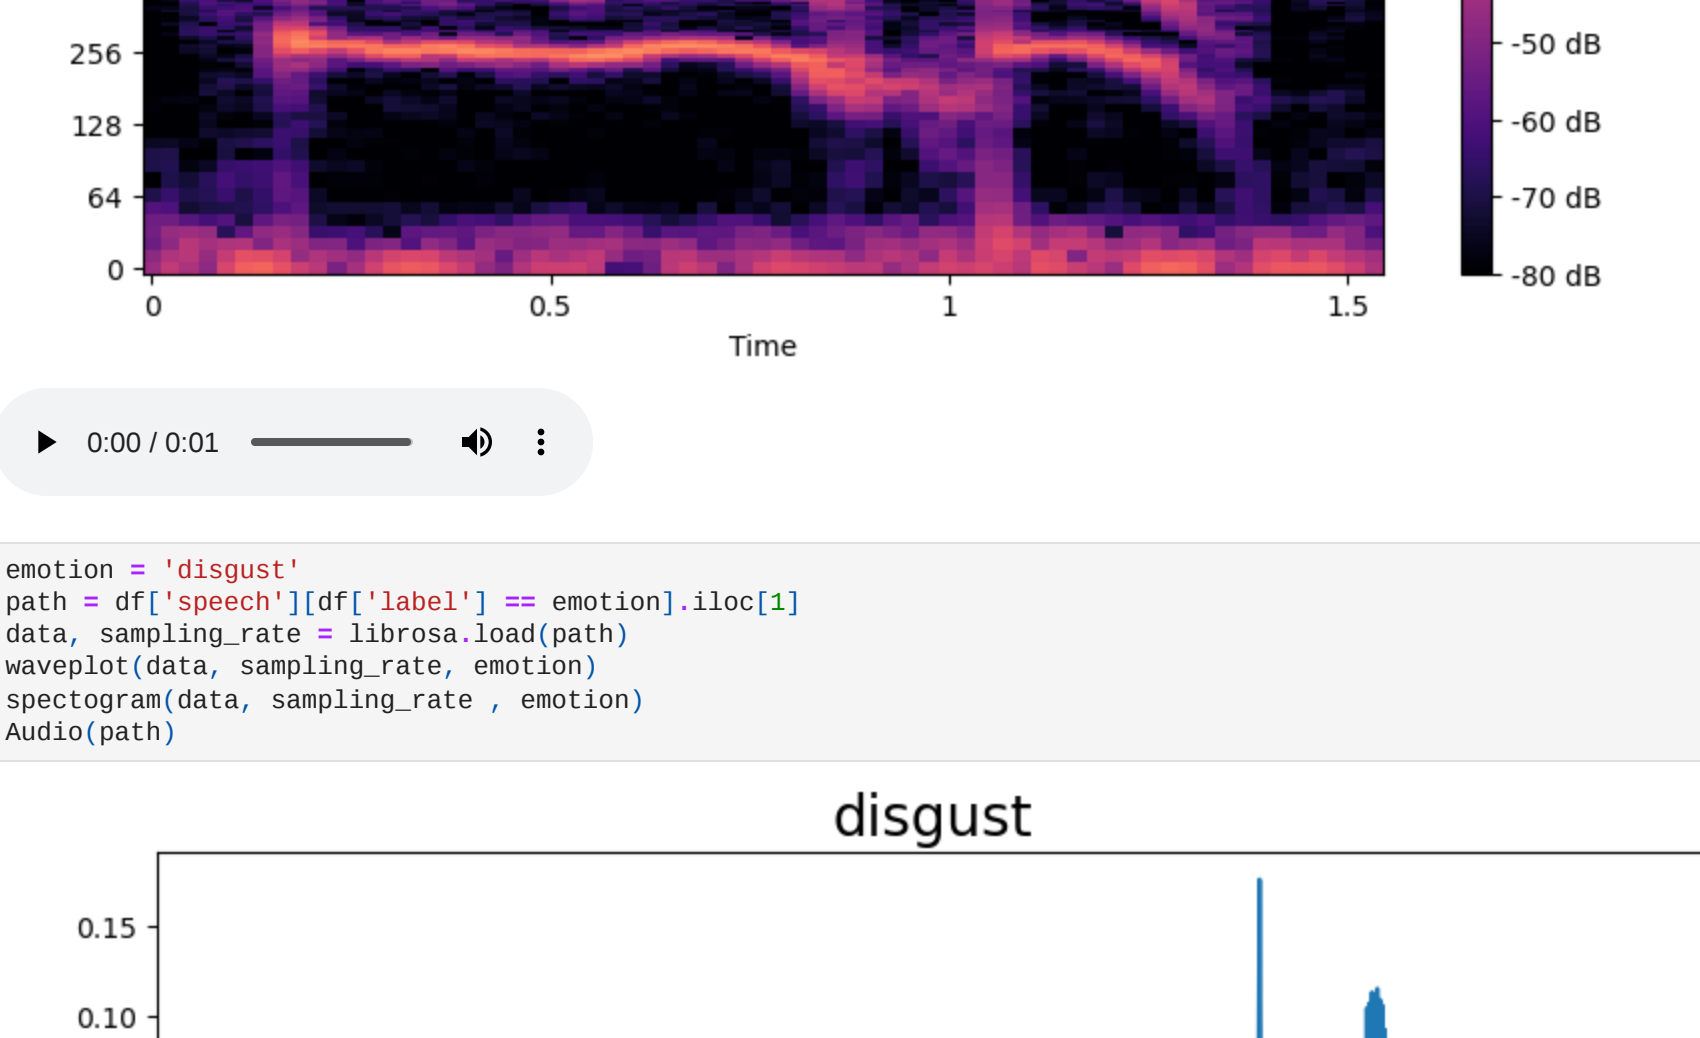
## Exploratory Data Analysis

```
In [9]: sns.countplot(data=df, x='label')
plt.xlabel('Emotion Label')
plt.ylabel('Count')
plt.title('Distribution of Emotions')
plt.xticks(rotation=90)
plt.show()
```



```
In [10]: def waveplot(data, sr, emotion):
plt.figure(figsize=(10, 4))
plt.title(emotion, size=20)
plt.plot(data)
plt.xlabel('Time (s)')
plt.ylabel('Amplitude')
plt.show()
def spectrogram(data, sr, emotion):
plt.figure(figsize=(10, 4))
plt.title(emotion, size=20)
S = librosa.amplitude_to_db(np.abs(librosa.stft(data)), ref=np.max)
librosa.display.specshow(S, sr=sr, x_axis='time', y_axis='log')
plt.colorbar(format='%+2.0f dB')
plt.show()
```

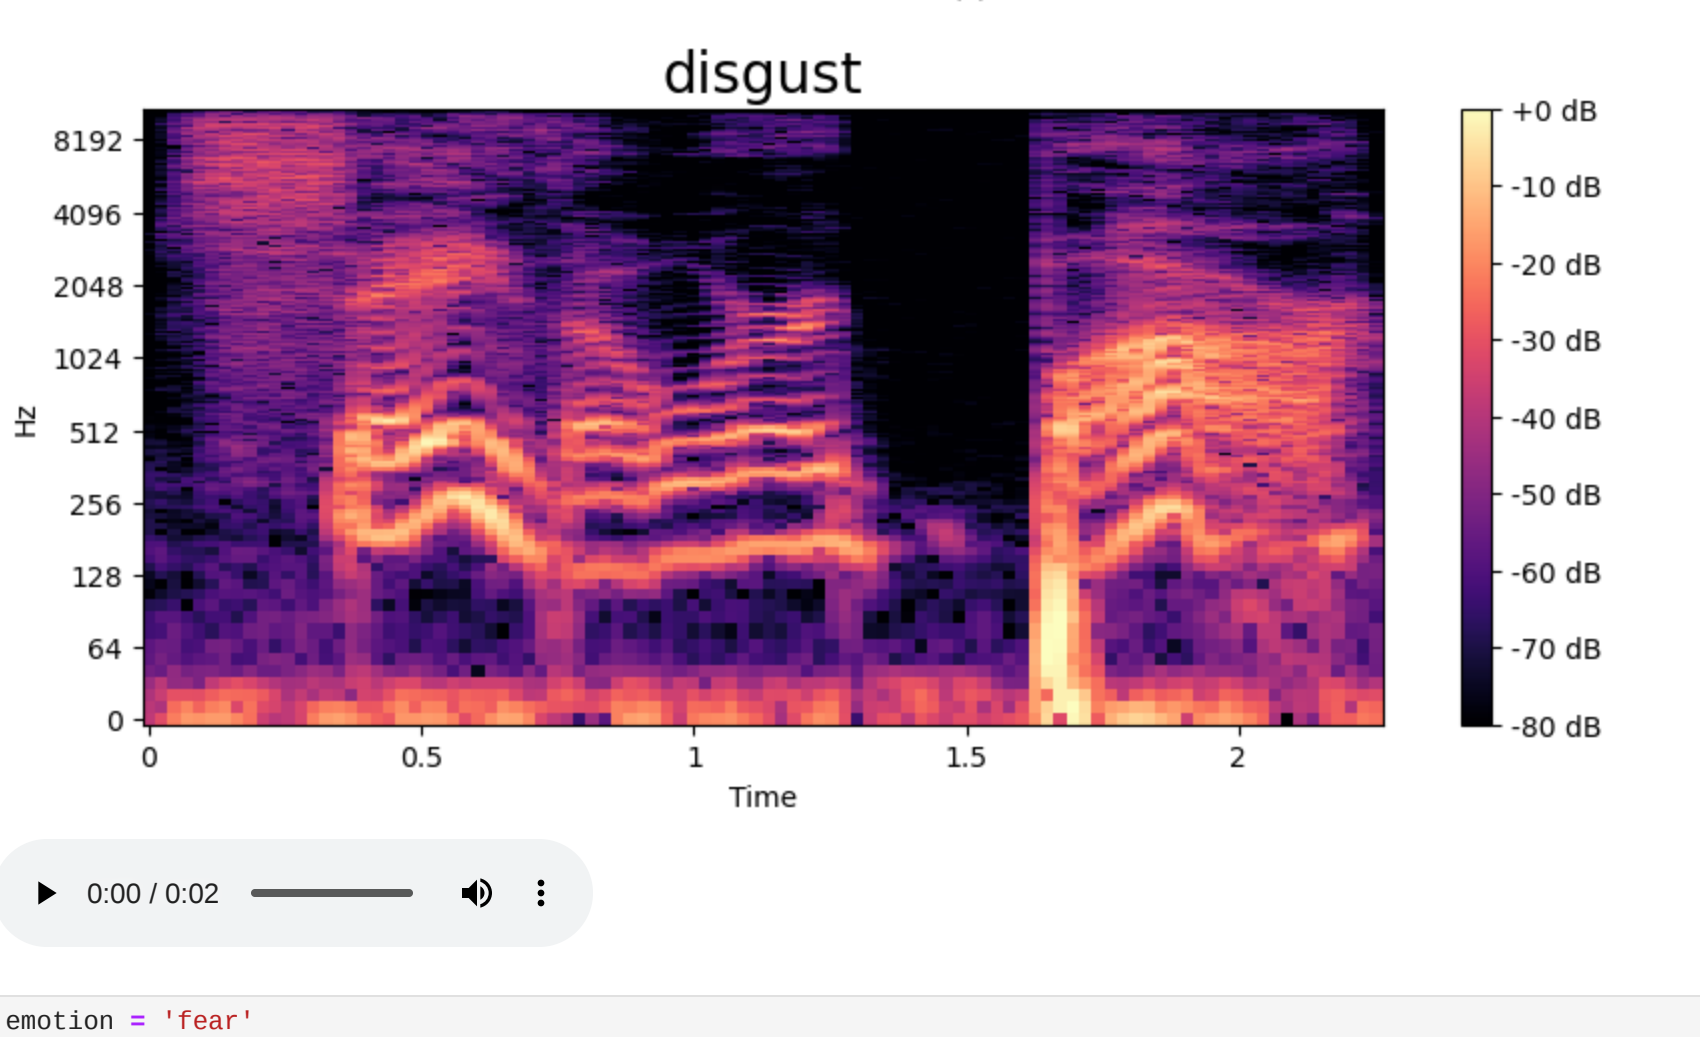
```
In [11]: emotion = 'angry'
path = df['speech'][df['label'] == emotion].iloc[0]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



```
Out[11]:
```

0:00 / 0:01

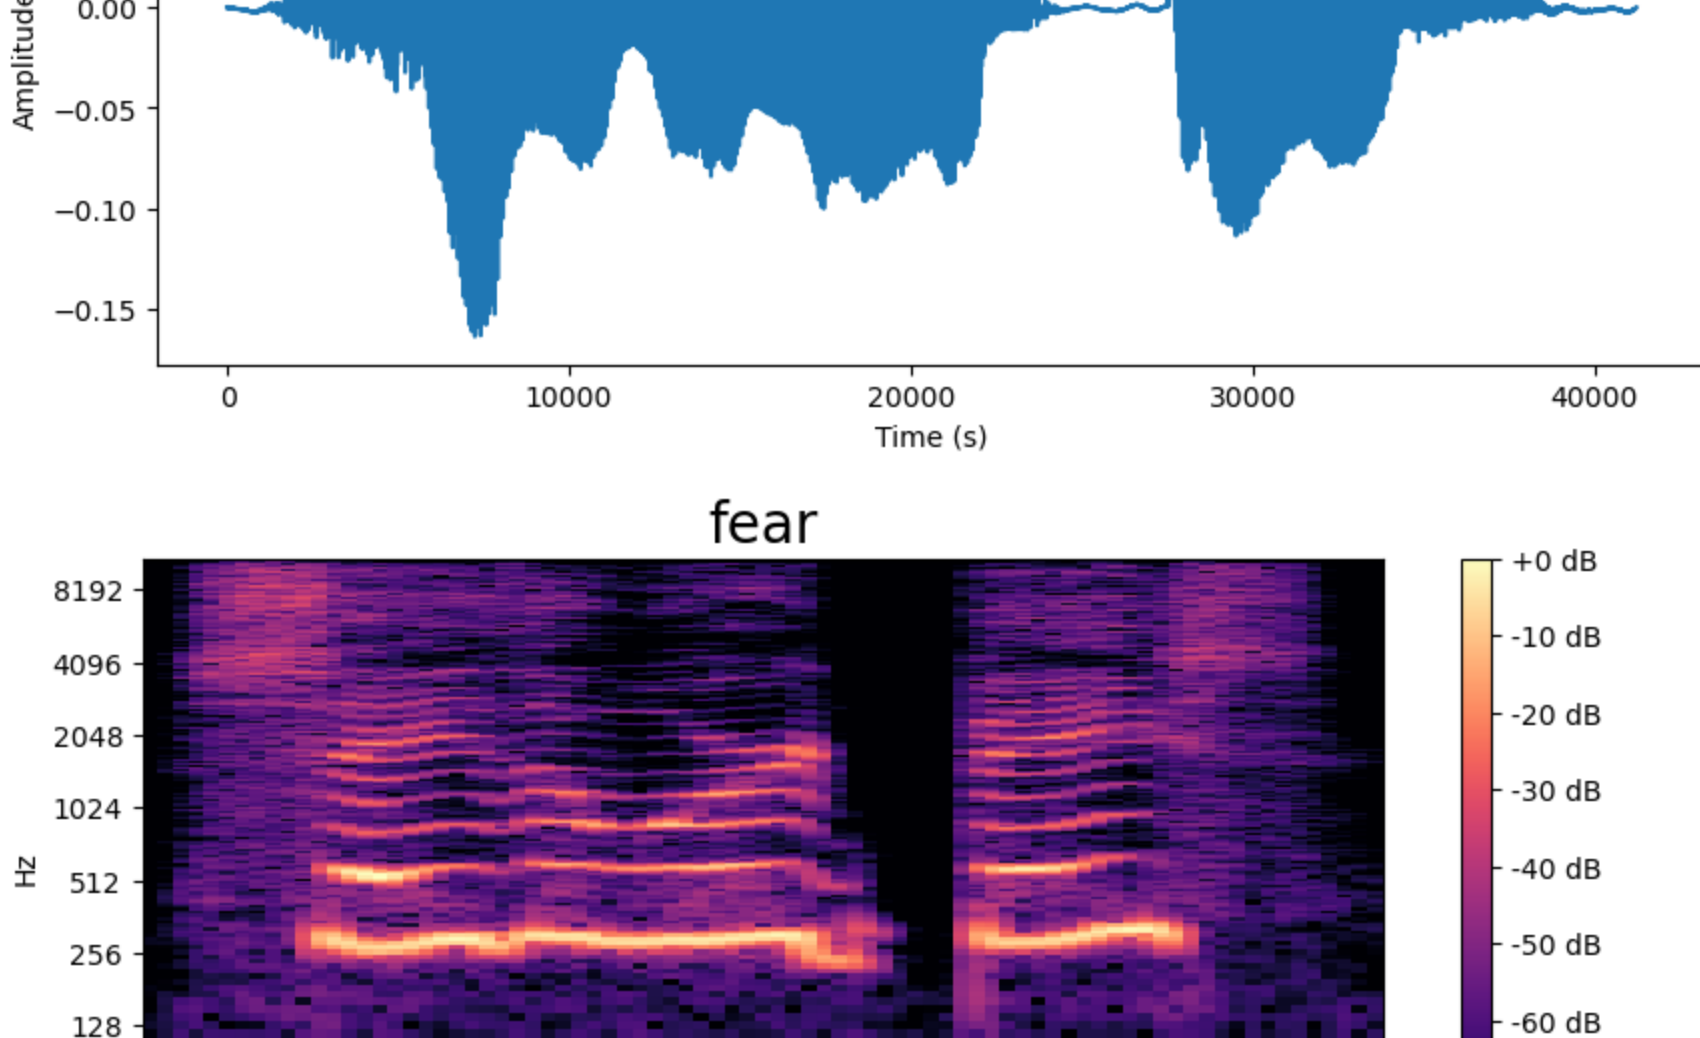
```
In [12]: emotion = 'disgust'
path = df['speech'][df['label'] == emotion].iloc[1]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



```
Out[12]:
```

0:00 / 0:02

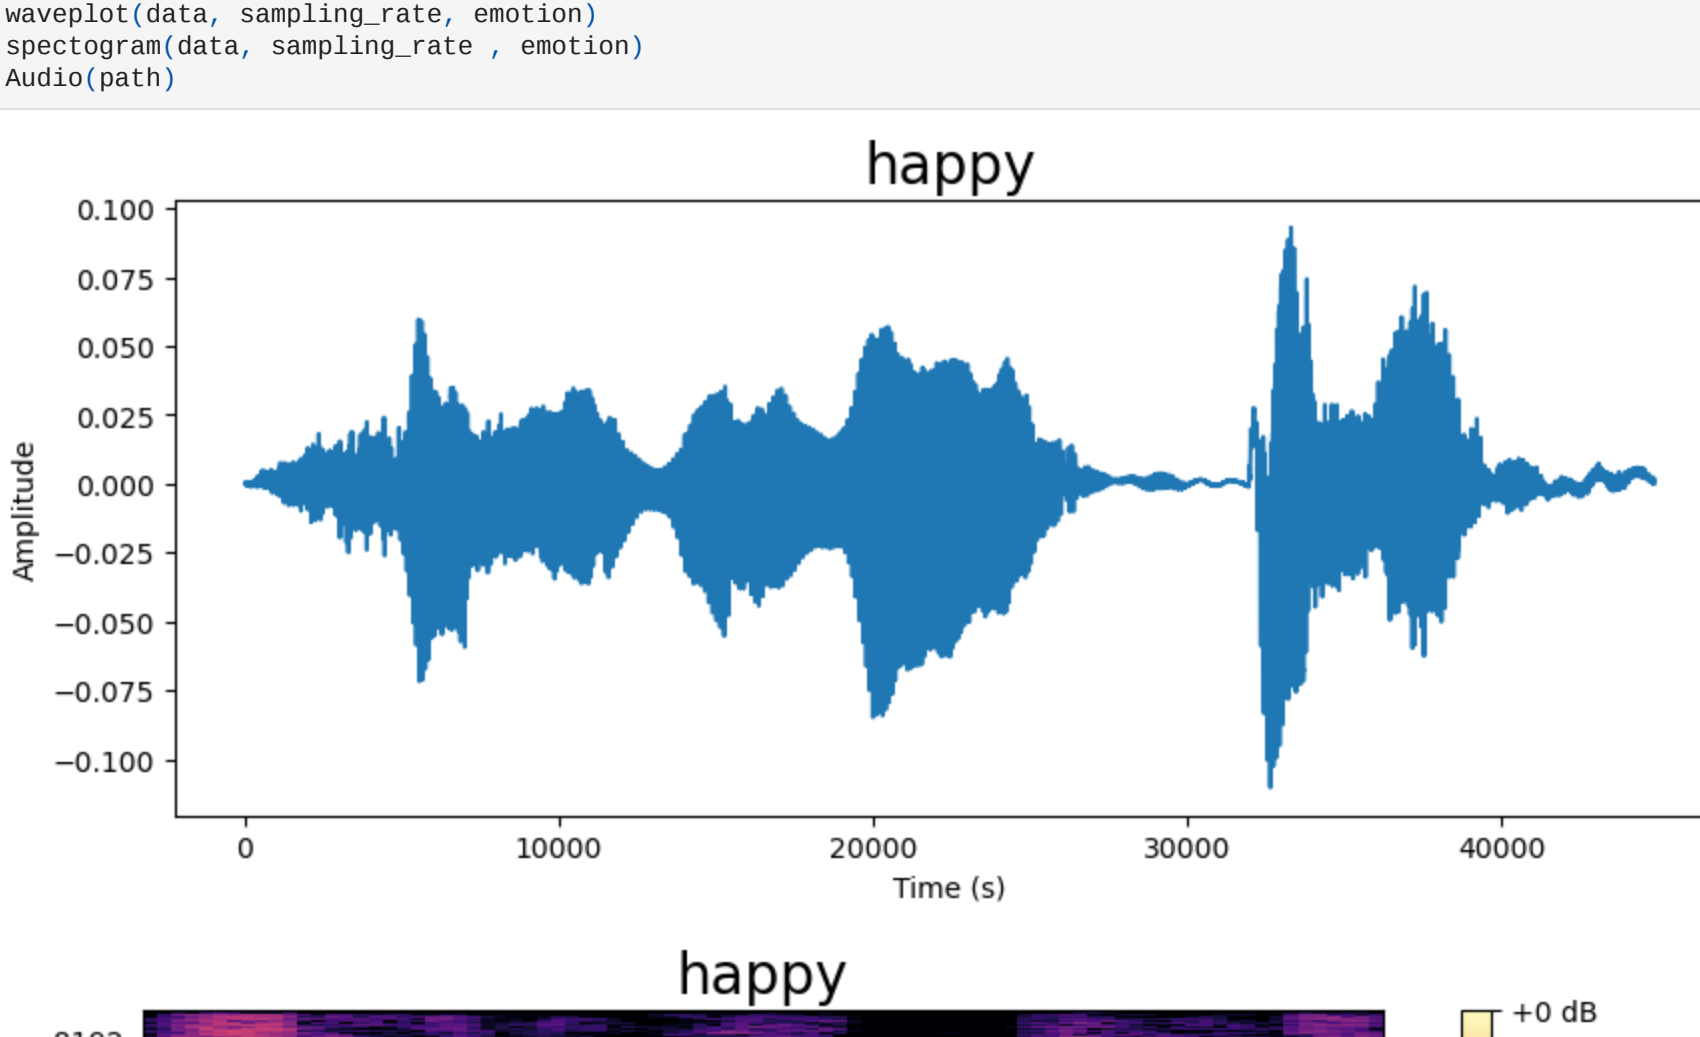
```
In [13]: emotion = 'fear'
path = df['speech'][df['label'] == emotion].iloc[2]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



```
Out[13]:
```

0:00 / 0:01

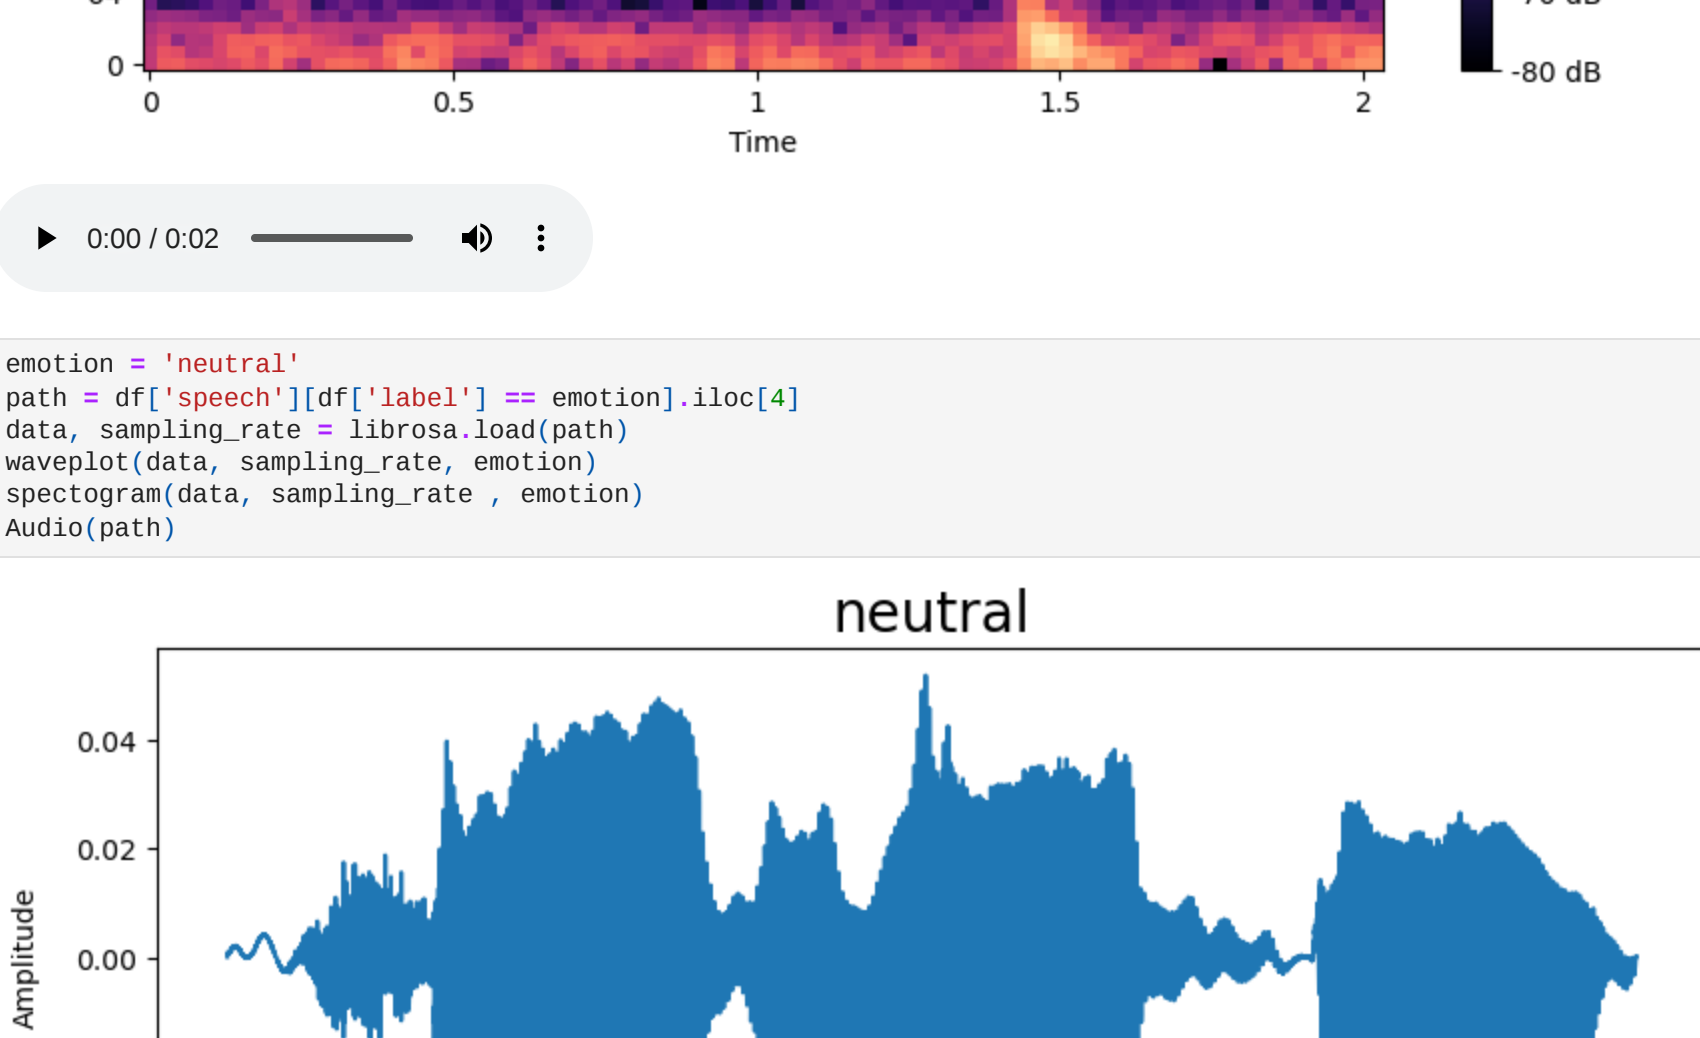
```
In [14]: emotion = 'happy'
path = df['speech'][df['label'] == emotion].iloc[3]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



```
Out[14]:
```

0:00 / 0:02

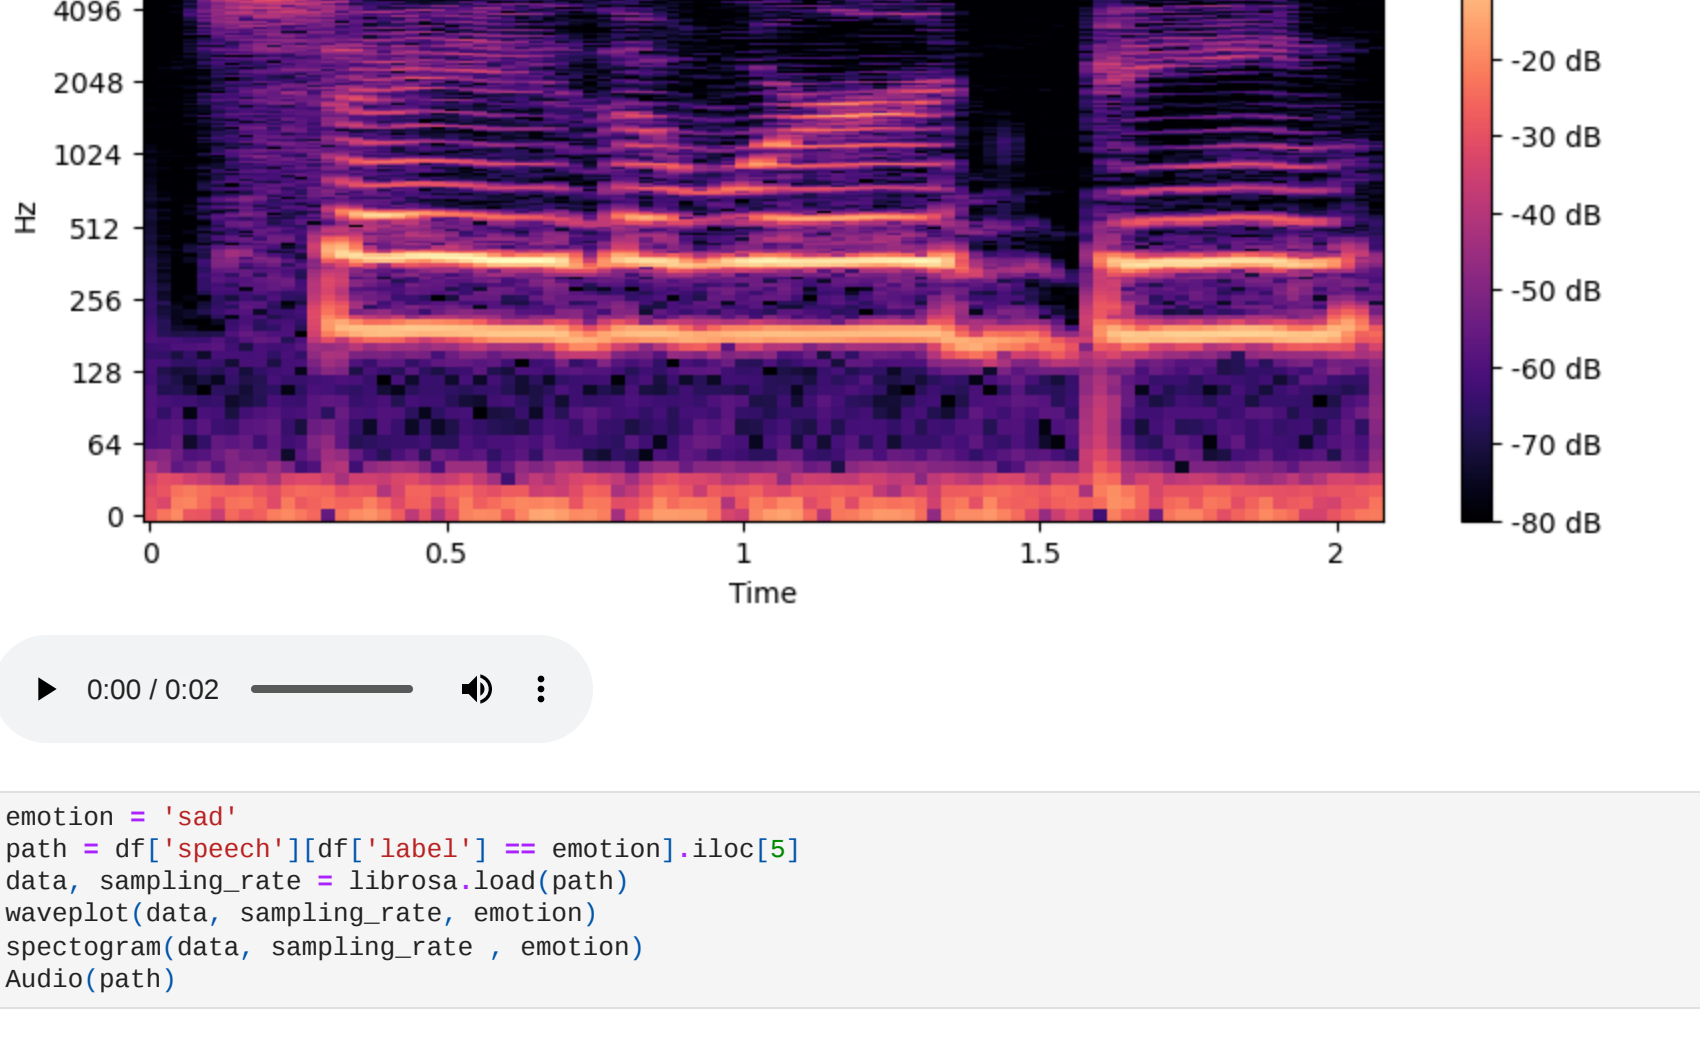
```
In [15]: emotion = 'neutral'
path = df['speech'][df['label'] == emotion].iloc[4]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



```
Out[15]:
```

0:00 / 0:02

```
In [16]: emotion = 'sad'
path = df['speech'][df['label'] == emotion].iloc[5]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



```
Out[16]:
```

0:00 / 0:02

## Feature Extraction

```
In [17]: def extract_mfcc(filename):
y, sr = librosa.load(filename, duration=3, offset=0.5)
mfcc = np.mean(librosa.feature.mfcc(y=y, sr=sr, n_mfcc=40).T, axis=0)
return mfcc
```

```
In [18]: extract_mfcc(df['speech'][0])
Out[18]: array([-3.9698628e+02, 7.7440535e+01, -1.9592790e+01, -2.1666689e+01,
-2.1127596e+08, 1.0975362e+01, -2.0360708e+01, -6.0924491e+08,
-7.2128259e+08, -5.5786940e+01, 2.8325327e+08, 2.0249514e+01,
7.2755128e+01, 1.3173770e+08, 2.8863375e+08, 2.8557933e+08,
-4.7129152e+08, -4.4351056e+08, 1.6211320e+08, -1.0239848e+08,
-5.5512615e+08, -1.7968802e+08, 7.0376532e+08, 6.4365647e+08,
6.3358549e+08, 2.1723694e+01, 1.9216989e+01, 2.0348926e+01,
1.3413365e+08, 8.3381745e+08, 3.9672240e+01, 0.11131430e+08,
9.5687389e+08, 5.4548682e+08, 2.5099637e+08, -1.8239714e+08,
1.8689627e+08, 9.3139238e+08, 2.0891501e+08, -1.906449120e+08],
dtype=float32)
```

```
In [19]: X_mfcc = df['speech'].apply(lambda x: extract_mfcc(x))
Out[19]: X_mfcc
```

```
Out[20]:
```

	X_mfcc
0	[ -396.9862, 77.44054, -19.59279, -21.66689, ...,
1	[ -465.73267, 98.73773, 0.8566806, 32.74544, ...
2	[ -429.79199, 46.124, 1.550476, 6.2170939, 2., ...
3	[ -401.46118, 76.32369, -12.531774, -25.288859, ...
4	[ -434.05756, 77.4405, 18.8655, 16.092943, 8.04., ...

```
5595 [ -486.48053, 80.379875, 32.462395, 46.578094, ...
5596 [ -425.90933, 102.54757, 24.800641, 42.044096, ...
5597 [ -378.50494, 80.81862, 35.398323, 39.74792, 2., ...
5598 [ -434.8814, 89.80284, 28.37329, 39.571071, -2., ...
5599 [ -421.8343, 70.89789, 32.479387, 45.642555, 4., ...
Name: speech, dtype: object
```

```
In [21]: X = [x for x in X_mfcc]
X = np.array(X)
X.shape
Out[21]: (5600, 40)
```

```
In [22]: X = np.expand_dims(X, -1)
X.shape
Out[22]: (5600, 40, 1)
```

```
In [23]: from sklearn.preprocessing import OneHotEncoder
enc = OneHotEncoder()
y = enc.fit_transform(df['label'])
```

```
In [24]: y = y.toarray()
y.shape
Out[25]: (5600, 7)
```

## Create the LSTM Model

```
In [26]: from keras.layers import Sequential
from keras.layers import Dense, LSTM, Dropout

model = Sequential([
    LSTM(256, return_sequences=False, input_shape=(40, 1)),
    Dropout(0.2),
    Dense(128, activation='relu'),
    Dropout(0.2),
    Dense(64, activation='relu'),
    Dropout(0.2),
    Dense(7, activation='softmax')
])

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()
```

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 256)	264192
dropout (Dropout)	(None, 256)	0
dense (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dropout_2 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 7)	455

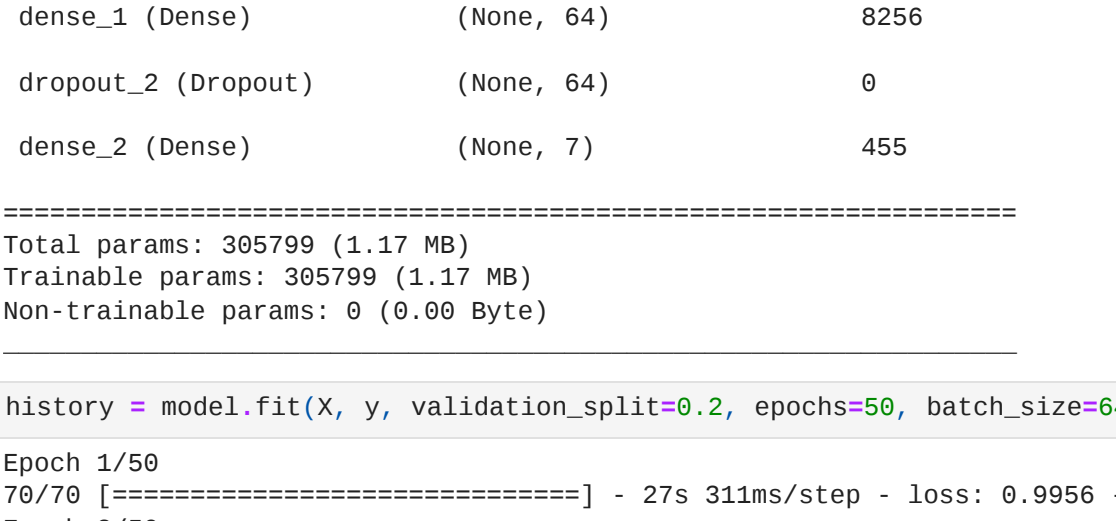
Total params: 305799 (1.17 MB)  
Trainable params: 305799 (1.17 MB)  
Non-trainable params: 0 (0.00 B)

```
In [27]: history = model.fit(X, y, validation_split=0.2, epochs=50, batch_size=64)
Epoch 1/50
76/78 [=====] - 0.75 318ms/step - loss: 0.9956 - accuracy: 0.6232 - val_loss: 0.2969 - val_accuracy: 0.9680
Epoch 2/50
76/78 [=====] - 0.21 280ms/step - loss: 0.2954 - accuracy: 0.8605 - val_loss: 0.1678 - val_accuracy: 0.9384
Epoch 3/50
76/78 [=====] - 0.20 289ms/step - loss: 0.1988 - accuracy: 0.8387 - val_loss: 0.0846 - val_accuracy: 0.9848
Epoch 4/50
76/78 [=====] - 0.21 299ms/step - loss: 0.1293 - accuracy: 0.8668 - val_loss: 0.0579 - val_accuracy: 0.9848
Epoch 5/50
76/78 [=====] - 0.21 288ms/step - loss: 0.1314 - accuracy: 0.8915 - val_loss: 0.0270 - val_accuracy: 0.9955
Epoch 6/50
76/78 [=====] - 0.19 268ms/step - loss: 0.0896 - accuracy: 0.8734 - val_loss: 0.0755 - val_accuracy: 0.9759
Epoch 7/50
76/78 [=====] - 0.19 276ms/step - loss: 0.0696 - accuracy: 0.8967 - val_loss: 0.0235 - val_accuracy: 0.9929
Epoch 8/50
76/78 [=====] - 0.19 276ms/step - loss: 0.0632 - accuracy: 0.8926 - val_loss: 0.0456 - val_accuracy: 0.9848
Epoch 9/50
76/78 [=====] - 0.19 269ms/step - loss: 0.0638 - accuracy: 0.9012 - val_loss: 0.0246 - val_accuracy: 0.9907
Epoch 10/50
76/78 [=====] - 0.19 276ms/step - loss: 0.0544 - accuracy: 0.8648 - val_loss: 0.0799 - val_accuracy: 0.9750
Epoch 11/50
76/78 [=====] - 0.20 287ms/step - loss: 0.0632 - accuracy: 0.8907 - val_loss: 0.0260 - val_accuracy: 0.9911
Epoch 12/50
76/78 [=====] - 0.19 268ms/step - loss: 0.0348 - accuracy: 0.8987 - val_loss: 0.0260 - val_accuracy: 0.9911
Epoch 13/50
76/78 [=====] - 0.19 273ms/step - loss: 0.0642 - accuracy: 0.8915 - val_loss: 0.0439 - val_accuracy: 0.9232
Epoch 14/50
76/78 [=====] - 0.20 279ms/step - loss: 0.0287 - accuracy: 0.9015 - val_loss: 0.0230 - val_accuracy: 0.9946
Epoch 15/50
76/78 [=====] - 0.20 279ms/step - loss: 0.0470 - accuracy: 0.8846 - val_loss: 0.0246 - val_accuracy: 0.9955
Epoch 16/50
76/78 [=====] - 0.19 282ms/step - loss: 0.0494 - accuracy: 0.8864 - val_loss: 0.0127 - val_accuracy: 0.9955
Epoch 17/50
76/78 [=====] - 0.19 273ms/step - loss: 0.0331 - accuracy: 0.8993 - val_loss: 0.0170 - val_accuracy: 0.9937
Epoch 18/50
76/78 [=====] - 0.20 287ms/step - loss: 0.0237 - accuracy: 0.8926 - val_loss: 0.0160 - val_accuracy: 0.9973
Epoch 19/50
76/78 [=====] - 0.20 290ms/step - loss: 0.0341 - accuracy: 0.8965 - val_loss: 0.0180 - val_accuracy: 0.9946
Epoch 20/50
76/78 [=====] - 0.21 301ms/step - loss: 0.0094 - accuracy: 0.9087 - val_loss: 0.0039 - val_accuracy: 0.9982
Epoch 21/50
76/78 [=====] - 0.21 301ms/step - loss: 0.0088 - accuracy: 0.9078 - val_loss: 0.0092 - val_accuracy: 0.9973
Epoch 22/50
76/78 [=====] - 0.20 289ms/step - loss: 0.0416 - accuracy: 0.8993 - val_loss: 0.0070 - val_accuracy: 0.9962
Epoch 23/50
76/78 [=====] - 0.20 281ms/step - loss: 0.0041 - accuracy: 0.8955 - val_loss: 0.0040 - val_accuracy: 0.9929
Epoch 24/50
76/78 [=====] - 0.21 299ms/step - loss: 0.0183 - accuracy: 0.8942 - val_loss: 0.0192 - val_accuracy: 0.9937
Epoch 25/50
76/78 [=====] - 0.21 284ms/step - loss: 0.0166 - accuracy: 0.8965 - val_loss: 0.0046 - val_accuracy: 0.9991
Epoch 26/50
76/78 [=====] - 0.19 277ms/step - loss: 0.0328 - accuracy: 0.8792 - val_loss: 0.0215 - val_accuracy: 0.9929
Epoch 27/50
76/78 [=====] - 0.19 276ms/step - loss: 0.0328 - accuracy: 0.8942 - val_loss: 0.0141 - val_accuracy: 0.9937
Epoch 28/50
76/78 [=====] - 0.22 313ms/step - loss: 0.0192 - accuracy: 0.8948 - val_loss: 0.0056 - val_accuracy: 0.9973
Epoch 29/50
76/78 [=====] - 0.21 305ms/step - loss: 0.0242 - accuracy: 0.8926 - val_loss: 0.0069 - val_accuracy: 0.9911
Epoch 30/50
76/78 [=====] - 0.21 289ms/step - loss: 0.0096 - accuracy: 0.8980 - val_loss: 0.0030 - val_accuracy: 0.9991
Epoch 31/50
76/78 [=====] - 0.21 296ms/step - loss: 0.0123 - accuracy: 0.8969 - val_loss: 0.0028 - val_accuracy: 0.9991
Epoch 32/50
76/78 [=====] - 0.20 282ms/step - loss: 0.0061 - accuracy: 0.8984 - val_loss: 0.0025 - val_accuracy: 0.9982
Epoch 33/50
76/78 [=====] - 0.20 289ms/step - loss: 0.0039 - accuracy: 0.8997 - val_loss: 0.0023 - val_accuracy: 0.9982
Epoch 34/50
76/78 [=====] - 0.19 274ms/step - loss: 0.0372 - accuracy: 0.8993 - val_loss: 0.0021 - val_accuracy: 0.9991
Epoch 35/50
76/78 [=====] - 0.20 288ms/step - loss: 0.0372 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 36/50
76/78 [=====] - 0.20 275ms/step - loss: 0.0076 - accuracy: 0.8997 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 37/50
76/78 [=====] - 0.18 269ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 38/50
76/78 [=====] - 0.18 269ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 39/50
76/78 [=====] - 0.18 259ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 40/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 41/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 42/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 43/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 44/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 45/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 46/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 47/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 48/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 49/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
Epoch 50/50
76/78 [=====] - 0.18 262ms/step - loss: 0.0076 - accuracy: 0.8999 - val_loss: 0.0026 - val_accuracy: 0.9991
```

## Plot the Results

```
In [28]: epochs = list(range(50))
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, label='train accuracy')
plt.plot(epochs, val_acc, label='val accuracy')
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend()
plt.show()
```



```
In [29]: loss = history.history['loss']
val_loss = history.history['val_loss']

plt.plot(epochs, loss, label='train loss')
plt.plot(epochs, val_loss, label='val loss')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend()
plt.show()
```

